

Flood - User Guide

Kajetan Rzepecki

September 25, 2013

Contents

1	Introduction	3
1.1	Use cases	3
1.2	Supported Protocols	3
1.3	Dependencies	3
1.4	Running Flood	3
2	So what's going on in here?	4
2.1	Simulated Users	4
2.2	User sessions	4
2.2.1	Session selection	4
2.2.2	Session inheritance	4
2.2.3	Actions & Event handlers	4
2.2.4	Timers & Counters	5
2.3	Flood phases	5
3	Test scenarios	6
3.1	Scenario file	6
3.2	Server setup	6
3.3	Phases setup	6
3.4	User session setup	8
3.5	User actions	9
3.6	Metadata	13
3.7	Example scenarios	14
3.7.1	Session inheritance	14
3.7.2	Ping-Pong	15
3.7.3	More examples	18
4	Test results & goals	19
4.1	Results format	19
4.2	Goal schemas	20
4.3	Continuous Integration integration	21
4.4	Example results	21
4.4.1	Session selection	21
4.4.2	Ping-Pong	22
4.4.3	More example results	23

1 Introduction

Flood is a load simulator useful for automatic **Comet/PUSH application** stress-testing. It is **asynchronous, event based** and enables you to create JSON encoded **test scenarios of arbitrary complexity** involving **tens of thousands of simulated users**, no Erlang required!

1.1 Use cases

Some of the most common use cases that **Flood** might be helpful in testing are:

- **Massive, real-time, on-line chats,**
- **Publisher-Subscriber channels,**
- **Instant messaging.**

However, Flood is general enough to test *any* event-based Comet application that uses the supported protocols.

1.2 Supported Protocols

Flood currently supports the **Socket.IO** protocol over **WebSocket** and **XHR-polling** transports with emphasis on Socket.IO event based communication. Flood also has *some* capabilities of using **raw HTTP** requests.

1.3 Dependencies

Flood uses several awesome libraries that are listed below. Since Flood is currently in development, *no particular stable versions are required* and by default the newest available versions are pulled in.

- Lager - a logging framework, found [here](#).
- Ibrowse - an HTTP client, found [here](#).
- Folsom - a metrics system, found [here](#).
- JSONx - a fast JSON parser, found [here](#).
- Jesse - a JSON Schema validator, found [here](#).
- `websocket_client` - a WebSocket client, found [here](#).

1.4 Running Flood

To run **Flood** simply run the **start-dev.sh** script and pass it the name of the Flood scenario to run:

```
$ ./start-dev.sh path/to/test.json
```

More about the format of the test files and Flood in general can be found in the following sections.

2 So what's going on in here?

This section describes what happens behind the scenes in **Flood** and how it reflects its usage. For the description of Flood test scenario files check [here](#). For the description of Flood test result files check [here](#).

Flood is a load simulator and its main purpose is **spawning a lot of simulated users** which perform certain, predefined **actions** (organized into **sessions**) and checking if certain **goals** (organized into **phases**) have been reached. Based on the goals Flood determines whether a test passes or fails. The big idea here is that all the tests and checks are performed automatically on a *large scale*.

2.1 Simulated Users

The core of the Flood are its simulated users. A user is a Finite State Machine currently consisting of three states:

- **disconnected** - user was spawned and initiated server connection and is awaiting a server response,
- **connected** - user received a server response and finalized all required handshakes and is performing actions defined in his session,
- **terminated** - the user is terminated and his process is removed. This is the final state.

All users start in **disconnected** state and attempt to connect to the server. Once **connected**, users start performing various actions and handling various events defined in their user sessions. In case of any unexpected connection problems, users fall back to the **disconnected** state and attempt to reconnect. Otherwise, users handle incoming events until they are explicitly terminated, either by a certain actions outcome or when the test is over.

User processes are lightweight and isolated, so tens of thousands of simulated users can exist simultaneously without any problems.

2.2 User sessions

User sessions are the meat of the Flood scenarios. Sessions describe the behaviour a user should exhibit after connecting to the server and are used to approximate real users behaviour during the simulation. A single scenario may define multiple user sessions that model very different behaviour. Each simulated user can only follow **one** session during its life span, but one session can be followed by multiple users distributed among several test phases.

2.2.1 Session selection

Sessions are **weighted**, meaning that each user session description declares a weight that corresponds to the probability of choosing that particular session over other sessions.

There are times where user sessions for particular users have to be selected from a set of multiple session descriptions. In such cases, a fitness proportionate selection algorithm determines the concrete user-session pairing.

2.2.2 Session inheritance

Some sessions may extend multiple other sessions. In such cases, the definitions of actions to be performed by users following a session that extends other sessions are combined with the definitions of actions of the inherited/extended/base sessions. This is very reminiscent of how **Common Lisp Object System's** multiple class inheritance works and the exact details of the implementation are omitted in this user guide. For a quick briefing check this [guide](#).

In case of multiple session inheritance, the exact ordering of actions to be performed is determined by Top-Sorting the session inheritance graph, and ensuring that the relative order of sessions at each inheritance level is preserved.

2.2.3 Actions & Event handlers

Sessions consist of **user actions** - *atomic* operations such as emitting a reply, incrementing a counter or starting a timer for future use. User session descriptions may define multiple **ordered** actions which will be performed by the simulated users after the server connection is established.

Some of the actions might be associated with several different **event handlers** that will execute them whenever a specific event is triggered. There are currently three types of event handlers supported:

- **Socket.IO message handlers** - triggered whenever a Socket.IO messages is received. These event handlers dispatch on the **opcode** of the received Socket.IO message.
- **Socket.IO event handlers** - most *unfortunately* named handlers that are triggered whenever a Socket.IO message with the **event** opcode is received. These event handlers dispatch on the **name of the Socket.IO event**.
- **timer timeout handlers** - triggered whenever a named timer started by the simulated user is due. These event handlers dispatch on the **name of the timer** that is timeouting.

In summary, actions are performed at user start-up and later whenever a handled event is received, be it a specific Socket.IO message, a Socket.IO event or a timer timeout. Some actions may cause more actions to be performed or even new event handlers to be created. Available actions and their semantics are described in a later section.

2.2.4 Timers & Counters

The last core concept in Flood are **counters** and **timers** which can be explicitly managed by the simulated users.

Counters are simple integers shared between all simulated users throughout a Flood test that can be **incremented**, **decremented** or **set** to a concrete value. Counters can be used to measure various quantities such as the total number of encountered errors, the number of messages sent/received etc. and can be later checked to ensure that certain thresholds have been reached (or not, in case of the errors).

Timers on the other hand are a little more complex. They can be **started**, **stopped** and **restarted**. Timers are **bounded** meaning that they will timeout after a certain time, resulting in a timeout event being generated and (ideally) handled. This makes defining some very complex behaviour possible. Similarly to counters, timers are used to measure a certain quantity - the time that passed between a timer start and a corresponding stop/timeout and can be used later in determining the tests outcome.

2.3 Flood phases

Phases are used to group simulated users and user sessions into logically distinct... Well, phases. Each Flood test is divided into several phases, each of which is scheduled and run at its own pace and with its own completion **goals** set. They run until said goals have been reached, or until they timeout.

Each phase defines how many simulated users it will support and which user sessions they will use. Phases are responsible for spawning users and periodically checking, whether a certain goal has been reached. Goals, being completely arbitrary **assertions on the values of counters and timers**, make it possible to determine whether a phase (and ultimately the entire Flood test) was succeeded or failed.

3 Test scenarios

This section describes the Flood scenario files and gives some general guidelines for writing them. Example scenarios can be found [here](#).

3.1 Scenario file

Flood uses JSON to encode test scenarios, no Erlang is required. Each scenario resides in a separate file and optionally several goal files (described in detail later). The overall structure of a Flood scenario consists of three required sections:

```
{
  "server" : {
    // Server setup.
  },

  "phases" : {
    // Test phases & goals.

    "phase_I" : {
      ...
    },
    ...
  },

  "sessions" : {
    // User session descriptions.

    "session_A" : {
      ...
    },
    ...
  }
}
```

3.2 Server setup

The **server** section is rather straightforward; it is used to setup the server connection. It has to define several mandatory fields:

```
"server" : {
  "host" : "",          // The server host.
  "port" : 0,           // The server port.
  "endpoint" : "",      // Endpoint used to connect to.
  "metadata" : {}       // Server-wide metadata (optional).
}
```

Example server configuration that will cause Flood to connect to `http://localhost:80/socket.io/1/` and define some server-wide metadata (more on metadata can be found [here](#)):

```
"server" : {
  "host" : "localhost",
  "port" : 80,
  "endpoint" : "/socket.io/1/",
  "metadata" : {
    "foo" : "bar"
  }
}
```

3.3 Phases setup

The **phases** section may define several arbitrarily named Flood phases. The ordering does not matter, as each phase explicitly names its start time.

```

"phases" : {
  "A" : {
    // A's description.
  },

  "B" : {
    // B's description.
  },
  ...
}

```

Each phase description has to follow this format:

```

"phase_I" : {
  "users" : 0,           // Number of users spawned during this phase.
  "user_sessions" : [], // Sessions spawned users should follow.

  "start_time" : 0,      // Time (in milliseconds) at which to start this phase.
  "spawn_duration" : 0,  // Duration (in milliseconds) Flood should take to spawn the users.

  "goal" : {},           // Goal of this phase (optional).
  "test_interval" : 0,   // Interval (in milliseconds) of the goal checks (optional).
  "timeout" : 0,         // Timeout (in milliseconds) of this phase (optional).

  "metadata" : {}        // Phase-wide metadata (optional).
}

```

The meaning of each of the fields is as follows:

- **users** - an integer number of users spawned during this phase. It is **mandatory**.
- **user_sessions** - a array of Flood user session names; the concrete user session will be selected at **random according to a sessions weight** (more about this can be found here). It is **mandatory**.
- **start_time** - an integer value that names a point in time (**in milliseconds**), relative to the start of the Flood, at which a phase should be started. It is **mandatory**.
- **spawn_duration** - an integer value that tells Flood how much time (**in milliseconds**) it should take to spawn **users** number of users. Users are spawned uniformly throughout this duration. Keep in mind that for various performance related reasons Flood **may actually take longer** to spawn the users, however it will never take less time to do so. This field is **mandatory**.
- **goal** - either an arbitrary JSON term that is a description of the goal of this phase (more on goals can be found here) or a string containing a path to the file containing the goal description relative to scenario file. This field is **optional**; not defining it will result in no goal checking whatsoever.
- **test_interval** - an integer value that tells Flood at what intervals (**in milliseconds**) in should check whether the goal has been reached. It is **optional**; not defining it will result in a single check at the phase timeout.
- **timeout** - an integer value that names a point in time (**in milliseconds**), relative to the start of the Flood, at which a phase should be terminated if it is still running. It is **optional**.
- **metadata** - a JSON object defining some phase-wide metadata (more on metadata later). It is **optional**.

Example phases setup:

```

"phases" : {
  "phase_I" : {
    "metadata" : { },

    "users" : 1000,
    "user_sessions" : ["session_A", "session_B"],

    "start_time" : 1000,

```

```

    "spawn_duration" : 1000
  },

  "phase_II" : {
    "metadata" : { },

    "users" : 1000,
    "user_sessions" : ["session_C"],

    "start_time" : 2000,
    "spawn_duration" : 5000

    "goal" : "./goal.jsonschema",
    "test_interval" : 100,
    "timeout" : 10000
  }
}

```

This setup will schedule two Flood phases. The first phase, **phase_I**, will start at 1000 ms and spawn 1000 users following either **session_A** or **session_B** over 1000 ms duration. The second phase, **phase_II**, will start at 2000 ms and spawn 1000 users following **session_C** over 5000 ms duration. Additionally, a **phase_II** goal check will be scheduled every 100 ms starting at 2000 ms and running until the goal provided in “./goal.jsonschema” file is met or until the phase timeout, set at 10000 ms, is reached.

3.4 User session setup

The **sessions** section may define several arbitrarily named Flood user sessions. The ordering does not matter, as each session explicitly names its relations to other sessions.

```

"sessions" : {
  "session_A" : {
    // A's description.
  },

  "session_B" : {
    // B's description.
  },
  ...
}

```

Each session description has to follow this format:

```

"session_A" : {
  "extends" : [], // Array of sessions extended by this session (optional).

  "weight" : 0.0, // Weight of this session (optional).
  "transport" : "", // Socket.IO transport used by this session (optional).

  "metadata" : {}, // Session-wide metadata (optional).
  "do" : [] // Array of actions to be performed by the user (optional).
}

```

The meaning of each of the fields is as follows:

- **extends** - an array of session names that this session extends (more about session inheritance can be found here). It is **optional** and omitting it means that this session does not extend any other sessions.
- **weight** - a real number determining how often simulated users will choose this session over other sessions (more on session selection can be found here); it is completely relative and depends on the total weight of a subset of sessions considered at one point (for examples at a certain Flood phase’s startup). It is **optional** and defaults to **0.0**.
- **transport** - a string naming a Socket.IO compatible transport protocol. It should be either of **websocket** or **xhr-polling**, but in general it is **optional** and defaults to the empty string.

- **metadata** - a JSON object defining some session-wide metadata (more on metadata later). It is **optional**.
- **do** - an array of actions to be performed by the users following this session (more on actions & event handlers can be found here; a list of all available actions can be found in the next section). It is **optional** and defaults to the empty array.

3.5 User actions

Actions are performed by the simulated users after their initialization and whenever an event triggers an event handler (for example, a Socket.IO message is received or a timer is due). Actions **ordering does matter** as some actions change the state of the simulated users.

Actions are represented as short JSON arrays consisting of an **action_ID** and a JSON object listing actions arguments:

```
["action_ID", {
  "argument_1" : "value_1", // Argument ordering does not matter.
  "argument_2" : "value_2",
  ...
}]
```

For convenience, some actions define a shorter forms that mean exactly the same, for example:

```
["action_ID", "value_1", "value_2"] // Mind the arguments ordering.
```

The following list lists available actions, describes their effects and arguments, and gives an example invocation in both full and short forms:

- **inc** - increments a named counter either by 1 or by **Value**. Example usage:

```
["inc", "counter_name"]
["inc", "counter_name", Value]
["inc", {
  "name" : "counter_name",
  "value" : Value
}]
```

- **dec** - decrements a named counter either by 1 or by **Value**. Example usage:

```
["dec", "counter_name"]
["dec", "counter_name", Value]
["dec", {
  "name" : "counter_name",
  "value" : Value
}]
```

- **set** - sets a named counter to a given **Value**. Example usage:

```
["set", "counter_name", Value]
["set", {
  "name" : "counter_name",
  "value" : Value
}]
```

- **start_timer** - starts a named timer timeouting in **Timeout** milliseconds. Example usage:

```
["start_timer", "timer_name", Timeout]
["start_timer", {
  "name" : "timer_name",
  "time" : Timeout
}]
```

- **stop_timer** - stops a named timer preventing it from timing out and triggering an event dispatch. Example usage:

```
["stop_timer", "timer_name"]
["stop_timer", {
  "name" : "timer_name"
}]
```

- `restart_timer` - restarts a named timer. Essentially, performs `stop_timer` and `start_timer` is quick succession. Example usage:

```
["restart_timer", "timer_name", Timeout]
["restart_timer", {
  "name" : "timer_name",
  "time" : Timeout
}]
```

- `timed` - executes a set of actions while timing their execution time which it then stores in a named counter. Results in whatever the actions result in. Example usage:

```
["timed", {
  "name" : "counter_name",
  "do" : [
    Action,
    ...
  ]
}]
```

- `on_timeout` - adds several timeout handlers to the simulated users state. If a given timeout handler already exists, new actions are appended **after** the existing ones, meaning they will be executed after the existing actions. Example usage:

```
["on_timeout", {
  "timer_name_1" : [
    Action,
    ...
  ],
  ...
}]
```

- `on_event` - adds several event handlers to the simulated users state. If a given event handler already exists, new actions are appended **after** the existing ones, meaning they will be executed after the existing actions. Example usage:

```
["on_event", {
  "event_1" : [
    Action,
    ...
  ],
  ...
}]
```

- `on_socketio` - adds several messages handlers to the simulated users state. If a given message handler already exists, new actions are appended **after** the existing ones, meaning they will be executed after the existing actions. Example usage:

```
["on_socketio", {
  "opcode_1" : [
    Action,
    ...
  ],
  ...
}]
```

- `emit_event` - emits Event with Args as a Socket.IO message with the event opcode. Example usage:

```
["emit_event", {
  "name" : Event,
  "args" : Args
}]
```

- **emit_socketio** - emits a Socket.IO message to the given **Endpoint** with the given **Opcode** and **Payload**. Example usage:

```
["emit_socketio", {
  "opcode" : Opcode,
  "endpoint" : Endpoint,
  "data" : Payload
}]
```

- **emit_http** - emits a synchronous HTTP request with a given **Method**, **Body**, **Headers** and **Timeout** to a given **Url**. Afterwards, executes actions defined in **on_reply** or **on_error** when the requests succeeded or failed respectively. Additionally, the response status code, headers and body can be accessed via **reply.status**, **reply.headers** and **reply.body** metadata in the **on_reply** branch. Example usage:

```
["emit_http", {
  "url" : Url,
  "method" : Method,
  "body" : Body,
  "headers" : Headers,
  "timeout" : Timeout,

  "on_reply" : [
    Action,
    ...
  ],

  "on_error" : [
    Action,
    ...
  ]
}]
```

- **match** - performs either a JSON-based or RegExp-based pattern-matching operation on **Subject**. RegExp-based matching takes precedence over JSON-based matching. The results are stored in the simulated users metadata under **Name.#** (where **#** is the index of the match) for RegExp-based matching or under respective **\$names** for JSON-based matching. Afterwards, executes actions defined in either **on_match** or **on_nomatch** when the matching succeeds or fails respectively. Example usage:

```
["match", {
  "name" : Name,
  "subject" : Subject,
  "re" : "regexp",

  "on_match" : [
    Action,
    ...
  ],

  "on_nomatch" : [
    Action,
    ...
  ]
}]
```

```
["match", {
  "subject" : Subject,
  "json" : {
```

```

        "field_1" : "$value_1",
        "field_2" : "$value_2",
        ...
    },

    "on_match" : [
        Action,
        ...
    ],

    "on_nomatch" : [
        Action,
        ...
    ]
}
}]

```

- **case** - performs a value case dispatch on a given **Value** selecting a matching **Branch** and executing its respective actions. Example usage:

```

["case", Value, {
    Branch : [
        Action,
        ...
    ],
    ...
}]
["case", {
    "condition" : Value,
    "branches" : {
        Branch : [
            Action,
            ...
        ],
        ...
    }
}]

```

- **def** - adds new metadata to the simulated users state. Example usage:

```

["def", {
    "key_1" : "value_1",
    "key_2" : "value_2",
    ...
}]

```

- **terminate** - immediately stops actions execution and terminates the simulated user with termination reason set to **Reason**. Disconnects him from the server and terminates his process. Example usage:

```

["terminate", Reason]
["terminate", {
    "reason" : Reason
}]

```

- **log** - prints a log line to the console formatting it with the **Format** and **Values**. The **Format** format is the same as Erlangs `io:format/2` (why yes, I did lie about the “no Erlang required” thing, deal with it). Example usage:

```

["log", Format, Values]
["log", {
    "format" : Format,
    "values" : Values
}]

```

- `!log` - a convenience action that allows easy `log` toggling; does nothing. Example usage:

```
[ "!log", Format, Values]
[ "!log", {
    "format" : Format,
    "values" : Values
}]
```

3.6 Metadata

Flood provides a per-user key-value store that can be accessed later by the simulated users. Various parts of a Flood scenario may define arbitrary key-value pairs in the `metadata` field. For example:

```
"metadata" : {
    "foo" : "bar",
    "bar" : [1, 2, 3],
    ...
}
```

Metadata defined in different sections has different scope. The `server` metadata is accessible by all the users. The `phase` metadata is accessible by the users spawned in that particular phase and `session` metadata is accessible by all the users following that metadata.

Metadata is **not shared** between users, instead every user accesses a unique copy. That means that the metadata can be freely modified added and removed during simulated users execution. This is the so-called *run-time metadata*.

Metadata from different sections **can and will shadow** metadata from other sections, the order is as follows (accessed from left to right):

```
run-time metadata >> session metadata >> phase metadata >> server metadata
```

Metadata can be accessed freely using *JSON Substitutions*:

```
[ "emit_event", {
    "name" : "$foo", // $foo --> "bar"
    "args" : "$bar"  // $bar --> [1, 2, 3]
}]
```

In general, JSON Substitutions can be used anywhere in the value position with the exception of **arrays of actions**, which are not substituted because they may contain their own Substitutions:

```
"do" : [
    "$some_action", // Not substituted.
    ["start_timer", "$timer", 1000] // Will be substituted when start_timer is executed.
]
```

There is some metadata that is added to the user state by default. Most of these correspond directly to the setup of different scenario sections:

- `server.host` - the server host,
- `server.port` - the server port,
- `server.endpoint` - the server endpoint,
- `server.url` - the server URL (host:port/endpoint),
- `server.sid` - the Socket.IO session ID received from the server,
- `server.heartbeat_timeout` - the Socket.IO heartbeat timeout received from the server,
- `server.reconnect_timeout` - the Socket.IO reconnect timeout received from the server,
- `server.available_transports` - the Socket.IO transports supported by the server,
- `phase.name` - the name of the *phase* the user was spawned in,

- `phase.users` - the number of users spawned in this *phase*,
- `phase.user_sessions` - the user sessions used in this *phase*,
- `phase.start_time` - the start time of this *phase*,
- `phase.spawn_duration` - the user spawn duration of this *phase*,
- `phase.test_interval` - the goal check interval of this *phase*,
- `phase.timeout` - the timeout time of this *phase*,
- `phase.goal` - the goal of this *phase*,
- `session.name` - the name of the *session* the user is following,
- `session.base_sessions` - the array of sessions extended by this *session*,
- `session.transport` - the Socket.IO transport used by this *session*,
- `session.weight` - the weight of this *session*.

Additionally, some temporary metadada may be added at various points to the user state. For example:

- `timer` - added when handling a timer timeout, contains the name of the timeouting timer,
- `event` - added when handling a Socket.IO event, contains the raw representation of the event,
- `event.name` - added when handling a Socket.IO event, contains the **name** of the event,
- `event.args` - added when handling a Socket.IO event, contains the **args** of the event,
- `message` - added when handling a Socket.IO message, contains the raw representation of the message,
- `message.opcode` - added when handling a Socket.IO message, contains the opcode of the message,
- `message.endpoint` - added when handling a Socket.IO message, contains the endpoint of the message,
- `message.data` - added when handling a Socket.IO message, contains the payload of the message.

3.7 Example scenarios

3.7.1 Session inheritance

This example shows session inheritance usage (more on this here). Full Flood scenario:

```
{
  "server" : {
    "host" : "localhost",
    "port" : 8080,
    "endpoint" : "/socket.io/1/"
  },

  "phases" : {
    "phase_I" : {
      "users" : 1,
      "user_sessions" : ["e"],

      "start_time" : 1000,
      "spawn_duration" : 1000,

      "timeout" : 3000
    }
  },

  "sessions" : {
    "a" : {
      "do" : [["log", "In A!"]]
    }
  }
}
```

```

    },

    "b" : {
      "extends" : ["a"],
      "do" : [["log", "In B!"]]
    },

    "c" : {
      "extends" : ["a"],
      "do" : [["log", "In C!"]]
    },

    "d" : {
      "extends" : ["b", "c"],
      "do" : [["log", "In D!"]]
    },

    "e" : {
      "weight" : 1.0,
      "transport" : "websocket",

      "extends" : ["d", "c", "b"],
      "do" : [["log", "In E!"]]
    }
  }
}

```

Sessions are composed retaining their topological ordering what ensures *sane* execution:

- session **e** extends **d**, **c** and **b** and requires them to run first **in order**,
- session **d** extends **b** and **c**,
- session **e** ensures that **b** and **c** *will* run, so **d** doesn't need to run **b** nor **c**,
- sessions **b** and **c** extend **a**,
- since session **d** requires both **b** and **c** to run and since **e** ensures that **b** and **c** *will* run, **d** only requires **a** to run first.

Flood output:

```

10:34:01.684 [notice] Running test examples/1.json
10:34:01.712 [notice] Scheduling Flood phase phase_I: 1 users every 1000 msecs (1 max)
                  starting at 1000 ms.
10:34:01.712 [notice] Scheduling Flood phase phase_I test at 3000 ms.
10:34:02.729 [notice] In A!
10:34:02.729 [notice] In D!
10:34:02.729 [notice] In C!
10:34:02.729 [notice] In B!
10:34:02.729 [notice] In E!
10:34:04.722 [notice] Flood phase phase_I reached its goal!

```

3.7.2 Ping-Pong

This example is a little more involved, it spawns 1000 users that ping a test server and measure the response time. It shows timers & counters usage (more on timers & counters here). Full Flood scenario:

```

{
  "server" : {
    "host" : "localhost",
    "port" : 8080,
    "endpoint" : "/socket.io/1/"
  },

```

```

"phases" : {
  "pingers" : {
    "users" : 1000,
    "user_sessions" : ["pinger"],

    "start_time" : 100,
    "spawn_duration" : 100,

    "test_interval" : 100,
    "timeout" : 10000,

    "goal" : {
      "type" : "object",
      "properties" : {
        "counters" : {
          "type" : "object",
          "properties" : {
            "received" : {
              "type" : "integer",
              "minimum" : 1000,
              "required" : true
            },
            "sent" : {
              "type" : "integer",
              "minimum" : 1000,
              "required" : true
            }
          }
        },
        "timers" : {
          "type" : "object"
        }
      }
    },

    "metadata" : {
      "ping_timeout" : 1000
    }
  },

  "sessions" : {
    "pinger" : {
      "transport" : "websocket",
      "weight" : 0.8,

      "do" : [
        ["on_socketio", {
          "1" : [
            ["log", "Ping ~s!", ["$server.sid"]],
            ["emit_event", {
              "name" : "ping",
              "args" : ["$server.sid"]
            }],
            ["inc", "sent"],
            ["start_timer", "ping", "$ping_timeout"]
          ],

          "5" : [
            ["inc", "received"],

```



```

        "95" : 463,
        "99" : 504,
        "999" : 523
    },
    "histogram" : {
        "x" : [124,184,244,304,364,454,554,654],
        "y" : [52,75,8,18,83,135,29,0]
    },
    "n" : 400
}
}
}

```

3.7.3 More examples

More Flood scenario examples and their results can be found in the **examples** directory of the Flood repository.

4 Test results & goals

This section describes the Flood test results and gives some general guidelines for interpreting them. Example results can be found [here](#).

4.1 Results format

Flood results are represented as JSON objects consisting of two main sections - **counters** containing final counter values and **timers** containing statistical analysis of the timers. The structure of the results file is as follows:

```
{
  "counters" {
    "counter_1" : 0, // Always a single value.
    ...
  },

  "timers" : {
    "timer_1" : {
      // Timer statistics.
    },
    ...
  }
}
```

Counters are **always** integers representing their **final value**. If a counter isn't used throughout the test (for example, an event triggering a counters increment is not received) it won't appear in the output of the Flood test.

Timers are more complicated as they have some statistical analysis done to them. They are represented as JSON objects of the following format:

```
"timer_1" : {
  "min" : 0,           // Minimum value recorded.
  "max" : 0,           // Maximum value recorded.
  "arithmetic_mean" : 0.0, // Arithmetic mean of samples.
  "geometric_mean" : 0.0, // Geometric mean of samples.
  "harmonic_mean" : 0.0,  // Harmonic mean of samples.
  "median" : 0,         // Median of samples.
  "variance" : 0.0,     // Variance of samples.
  "standard_deviation" : 0.0, // Standard deviation of samples.
  "skewness" : 0.0,     // Skewness of samples.
  "kurtosis" : 0.0,    // Kurtosis of samples.
  "percentile" : {
    "50" : 0,           // 50% percentile.
    "75" : 0,           // 75% percentile.
    "90" : 0,           // 90% percentile.
    "95" : 0,           // 95% percentile.
    "99" : 0,           // 99% percentile.
    "999" : 0           // 99.9% percentile.
  },
  "histogram" : {
    "x" : [0, ...],     // X axis values of the histogram (buckets).
    "y" : [0, ...]     // Y axis values of the histograms (samples).
  },
  "n" : 0               // The total number of samples.
}
```

The provided statistics are:

- **min** - the lowest sampled value,
- **max** - the highest sampled value,

- `arithmetic_mean` - a straightforward, arithmetic mean of the sampled values,
- `geometric_mean` - a less straightforward, geometric mean of the sampled values,
- `harmonic_mean` - a harmonic mean of the sampled values,
- `median` - the median of the sampled values,
- `standard_deviation` - the standard deviation of the sampled values,
- `variance` - the variance of the sampled values,
- `skewness` - the skeweness of the sampled values,
- `kurtosis` - the kurtosis of the sampled values,
- `percentile.50` - the 50% percentile of the sampled values, means that at least 50% of the samples are below or equal to this value,
- `percentile.75` - the 75% percentile of the sampled values, means that at least 75% of the samples are below or equal to this value,
- `percentile.90` - the 90% percentile of the sampled values, means that at least 90% of the samples are below or equal to this value,
- `percentile.95` - the 95% percentile of the sampled values, means that at least 95% of the samples are below or equal to this value,
- `percentile.99` - the 99% percentile of the sampled values, means that at least 99% of the samples are below or equal to this value,
- `percentile.999` - the 99.9% percentile of the sampled values, means that at least 99.9% of the samples are below or equal to this value,
- `hitogram.x` - the X axis values of a histogram of the sampled values (buckets).
- `hitogram.y` - the Y axis values of a histogram of the sampled values (samples).
- `n` - the number of samples used for the statistical analysis.

To properly interpret the results keep in mind that the samples are **collected within a 60 second sliding window** with **at most 100 uniformly selected samples collected every second**. This means that if there are many more timer updates per second, only 100 uniformly selected measurements will be averaged and added to the samples on which statistical analysis is performed. Furthermore, the values reflect the state of a timer in the past 60 seconds only and so global extremes may not appear in the result.

On the other hand, keep in mind that if there are too little samples available, no statistical analysis can and will be done, instead all values will default to 0 or won't be included in the output at all.

4.2 Goal schemas

Flood uses JSON schema compatible validator when testing whether goals have been reached or not. Every phase may specify a **goal** that has to be a JSON schema that will be used to check current values of the counters and timers (the format of the results JSON can be found in the previous section) or a **relative path** to a JSON schema file that should be used instead. For example:

```
"sample_phase_I" : {
  "goal" : {
    "type" : "object",
    "properties" : {
      "counters" : {
        "type" : "object",
        "properties" : {
          "counter_1" : {
            // JSON Schema to validate counter_1.
          },
          ...
        }
      }
    }
  }
}
```

```

    },
    "timers" : {
        "type" : "object",
        "properties" : {
            "timer_1" : {
                // JSON Schema to validate timer_1.
            },
            ...
        }
    }
}

},
...
},

"sample_phase_II" : {
    "goal" : "path/to/schema.jsonschema", // File containing goal schema.
    ...
}

```

Goals (if defined) are checked every `test_interval` milliseconds (if configured) or once at the phase `timeout` (if configured). If a goal check fails either nothing happens or another check is scheduled. On the other hand, if a goal check passes, a `phasename_goal_time` counter specifying the point in time (relative to the start of Flood) will be added to the named counters and later included in the results file (`phasename` part is the name of the respective phase).

If a `timeout` has been configured for any of the phases included in a scenario Flood will terminate as soon as the chronologically last timeout is reached, or when the last goal check passed, whichever comes first. Phases that end chronologically sooner will end and users spawned during their execution will be terminated.

The results file is dumped to the disk at Flood termination under `testname.flood.results.json` name (`testname` part is the base-name of the Flood scenario currently running).

4.3 Continuous Integration integration

Flood can be run automatically and easily integrated into any Continuous Integration environment. Flood will terminate with exit reason **1** or **0** when the test fails or succeeds respectively with logs saved in **log** directory and results dumped to the test scenario directory for future reference.

4.4 Example results

4.4.1 Session selection

Corresponds to the `examples/1.json` Flood scenario (note that there are no timers used in this test). Test goal:

```

{
    "type" : "object",
    "properties" : {
        "counters" : {
            "type" : "object",
            "properties" : {
                "xhr_clients" : {
                    "type" : "integer",
                    "minimum" : 180,
                    "maximum" : 220,
                    "required" : true
                },
                "websocket_clients" : {
                    "type" : "integer",
                    "minimum" : 780,
                    "maximum" : 820,
                    "required" : true
                }
            }
        }
    }
}

```

```

    }
  },
  "timers" : {
    "type" : "object"
  }
}
}

```

The result when the goal has been reached:

```

{
  "counters" : {
    "ws_incomming" : 780,
    "http_outgoing" : 1440,
    "ws_outgoing" : 0,
    "http_incomming" : 1220,
    "disconnected_users" : 0,
    "connected_users" : 1000,
    "alive_users" : 1000,
    "all_users" : 1000,
    "terminated_users" : 0,
    "xhr_clients" : 220,
    "sample_phase_goal_time" : 2800,
    "websocket_clients" : 780
  },
  "timers" : []
}

```

4.4.2 Ping-Pong

Corresponds to the **examples/2.json** Flood scenario. Test goal:

```

{
  "type" : "object",
  "properties" : {
    "counters" : {
      "type" : "object",
      "properties" : {
        "received" : {
          "type" : "integer",
          "minimum" : 1000,
          "required" : true
        },
        "sent" : {
          "type" : "integer",
          "minimum" : 1000,
          "required" : true
        }
      }
    }
  },
  "timers" : {
    "type" : "object"
  }
}
}

```

The result when the goal has been reached:

```

{
  "counters" : {
    "ws_incomming" : 2000,
    "http_outgoing" : 1000,

```

```

    "ws_outgoing" : 1000,
    "http_incomming" : 1000,
    "disconnected_users" : 0,
    "connected_users" : 1000,
    "pingers_goal_time" : 1900,
    "alive_users" : 1000,
    "all_users" : 1000,
    "terminated_users" : 0,
    "received" : 1000,
    "sent" : 1000
  },
  "timers" : {
    "ping" : {
      "min" : 54,
      "max" : 523,
      "arithmetic_mean" : 298.8575,
      "geometric_mean" : 260.985015508945,
      "harmonic_mean" : 216.292895973774,
      "median" : 347,
      "variance" : 17071.5510714286,
      "standard_deviation" : 130.658145828833,
      "skewness" : -0.387733104425692,
      "kurtosis" : -1.27787946255272,
      "percentile" : {
        "50" : 347,
        "75" : 401,
        "90" : 447,
        "95" : 463,
        "99" : 504,
        "999" : 523
      },
      "histogram" : {
        "x" : [124,184,244,304,364,454,554,654],
        "y" : [52,75,8,18,83,135,29,0]
      },
      "n" : 400
    }
  }
}

```

4.4.3 More example results

More Flood scenario examples and their results can be found in the `examples` directory of the Flood repository.